

ReAct Agent



ReAct = Reasoning + Acting

- It is a method for building agents where the LLM doesn't just give a direct answer, but:
 1. **Thinks step by step** (Reasoning)
 2. **Takes actions** using tools (Acting)
 3. **Uses results from tools** to continue reasoning
 4. **Finally answers the user**

Why is it needed?

- Normal LLMs can hallucinate (make up wrong answers).
- With **ReAct**, the agent can:
 - Call a **tool** (like Google, Wikipedia, Calculator, Database, API).
 - Check the tool's result.
 - Use that real information in the answer.

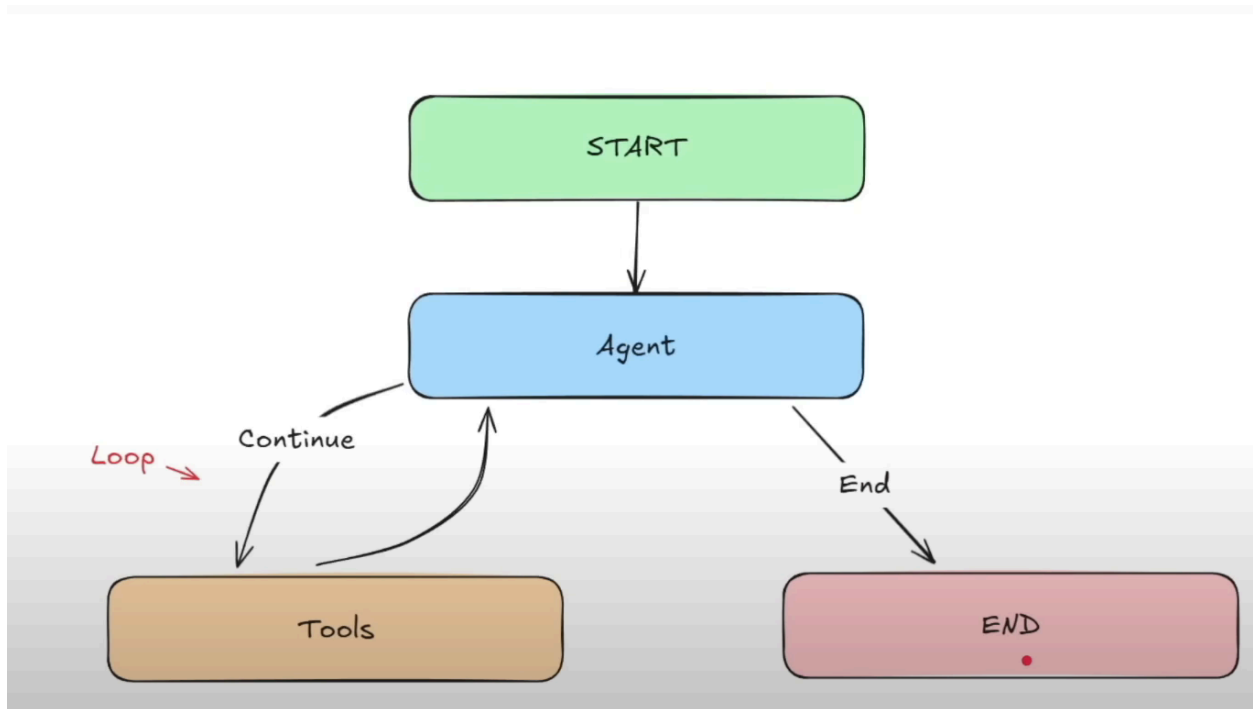
This makes it **more reliable**.

How does it work internally?

Imagine a conversation:

1. **User**: "What is the population of France divided by 2?"
2. **Agent (Reasoning)**: I don't know directly. Step 1: I should look up France's population.
3. **Agent (Action)**: Calls a tool (Wikipedia API).

4. **Tool Result:** "France population = 67 million"
5. **Agent (Reasoning):** Now divide 67 million by 2 = 33.5 million.
6. **Agent (Final Answer):** "It's about 33.5 million."



Code

```
from typing import Annotated, Sequence, TypedDict
from dotenv import load_dotenv
from langchain_core.messages import BaseMessage # The foundational class
for all message types in LangGraph
from langchain_core.messages import ToolMessage # Passes data back to LL
M after it calls a tool such as the content and the tool_call_id
from langchain_core.messages import SystemMessage
from langchain_groq import ChatGroq
from langchain_core.tools import tool
from langgraph.graph.message import add_messages
from langgraph.graph import StateGraph, END
from langgraph.prebuilt import ToolNode
```

```
load_dotenv()
```

Annotated → Allows you to **attach extra metadata** to a type hint. I

Sequence → It's used to say: "This function accepts any ordered collection, not just a list."

✓ Works with:

- `["a", "b", "c"]` (list)
- `("a", "b", "c")` (tuple)
- `"abc"` (string)

📌 Why use **Sequence** instead of **List** ?

- `List[str]` only accepts lists.
- `Sequence[str]` is **more flexible**, and works with any indexable iterable (list, tuple, str, etc.).

ToolMessage

- A special message used when a **tool has returned a result** and you're sending that result back to the LLM.
- Typically follows a **tool call** that the LLM made.

🧠 Purpose:

- Bridges the gap between the **LLM asking for a tool** and the **tool's actual response**.

✓ **BaseMessage**

- `BaseMessage` is the **abstract base class** for all message types in LangChain.
- Other message types (like `HumanMessage`, `AIMessage`, `SystemMessage`, `ToolMessage`, etc.) **inherit from** this.
- `BaseMessage` **is the parent class**



SystemMessage

- A message type used to pass **instructions, role-setting, or context** to the LLM.

add_messages

- This is a **reducer function**
- LangGraph agents typically maintain a **state** that contains a list of messages (like `HumanMessage`, `AIMessage`, etc.). `add_messages` makes it easier to **append new messages** to that list in a structured and consistent way.

```
from langgraph.graph.message import add_messages
from langchain_core.messages import AIMessage

def some_node(state):
    response = AIMessage(content="Sure, here's the info you asked for.")
    return add_messages(state, [response])
```

ToolMessage →

- A message type specifically used when an LLM calls a **tool**.
 - It contains the **result of a tool execution**.
 - Also carries the `tool_call_id` so LangChain knows which request this belongs to.

Example flow:

1. LLM says: *"Use calculator with input 2+2"*.
2. Tool is executed → result = `4`.
3. This result is wrapped in a `ToolMessage` and passed back to the LLM.

```
from langchain_core.tools import tool
```

`@tool` → Turns a function into a Tool

- This is a **decorator**. You put `@tool` above a Python function to magically convert it into a tool that the AI can recognize and call.

`ToolNode` →

A **ready-made LangGraph node** that handles **tool execution**.

- Instead of writing your own node to call a tool and process its result, you just use `ToolNode`.
- It takes care of:
 - Running the tool
 - Wrapping the result in `ToolMessage`
 - Sending it back to the LLM

A **pre-built node** that handles all the complexity of:

1. Receiving a list of tool calls from the AI.
2. Finding the correct tool.
3. Executing it.
4. Packaging the results into objects.

It saves you from writing this logic yourself.

```
class AgentState(TypedDict):  
    message: Annotated[Sequence[BaseMessage], add_messages]
```

`Sequence[BaseMessage]`

- `Sequence` = any ordered, indexable collection (like a `list` or `tuple`).
- `BaseMessage` = parent class for `HumanMessage`, `AIMessage`, `SystemMessage`, etc.

So, `"message"` will hold a **list of messages**:

```
[
  HumanMessage(content="Hello"),
  AIMessage(content="Hi, how can I help?")
]
```

Annotated[..., add_messages]

In LangGraph, that metadata (`add_messages`) tells the framework:

“When this node returns new messages, use the `add_messages` function to merge them into this field.”

The **special instruction** `add_messages` , LangGraph knows:

“Whenever I add new messages, I should not overwrite old ones, but append them to the list.”

Create Tool

```
@tool
def add(a: int, b:int):
    """This is an addition function that adds 2 numbers together"""

    return a + b

@tool
def subtract(a: int, b: int):
    """Subtraction function"""
    return a - b

@tool
def multiply(a: int, b: int):
    """Multiplication function"""
    return a * b
```



!!!If the doc string is removed, the function will not work.

Create a list of tools

```
#Create a list of tools
```

```
tools = [add, subtract, multiply]
```

Model + Tools:

```
#Model
```

```
model = ChatGroq(model="llama-3.3-70b-versatile").bind_tools(tools)
```

- In this way, model has access to the tools

Model Call Node:

```
def model_call(state: AgentState) → AgentState:  
    system_prompt = SystemMessage(content="You are my AI assistant, please answer my query to the best of your ability.")  
    response = model.invoke([system_prompt] + state["messages"])  
    return {"messages": [response]}
```

Output:

```
AIMessage(content="I'll do my best to provide a helpful and accurate response. What's your question?", additional_kwargs={}, response_metadata={'token_usage': {'completion_tokens': 19, 'prompt_tokens': 47, 'total_tokens': 66, 'completion_time': 0.046528529, 'prompt_time': 0.017985182, 'queue_time': 0.047918769, 'total_time': 0.064513711}, 'model_name': 'llama-3.3-70b-versatile', 'system_fingerprint': 'fp_2ddfb0da0', 'finish_reason': 'stop'})
```

```
p', 'logprobs': None}, id='run--3883e593-2eb1-4480-b079-6ea4c53ca203-0', usage_metadata={'input_tokens': 47, 'output_tokens': 19, 'total_tokens': 66})
```

`return {"messages": [response]}` → Compact way to update the state.



We are not returning a string because (`response.content`) → `"messages"` field **must hold a list of `BaseMessage` objects**, not plain strings.



We just write `"messages": [response]` because the `add_messages` function handles the appending of messages.

Conditional Edge

```
def should_continue(state: AgentState) → AgentState:
```

```
    messages = state["messages"]
```

```
    last_message = messages[-1]
```

```
    if not last_message.tool_calls:
```

```
        return "end"
```

```
    else:
```

```
        return "continue"
```

`messages[-1]` gets the **most recent message** (last one in the list)

What is `.tool_calls` ?

- Some messages (usually from the assistant) may contain **tool calls**.

- Example: the assistant might decide: *"I need to use the Wikipedia API tool"*.
- In that case, the message will have something like:

```
{
  "role": "assistant",
  "tool_calls": [{"id": "1", "name": "WikipediaAPIWrapper", "arguments": {"query": "cats"}}]
}
```

If `.tool_calls` is **empty**, it means the assistant is not calling any tool and is just giving a final response.

- If the **last message has no tool calls** → that means the assistant is done → return `"end"`.
- If the **last message contains tool calls** → that means the assistant wants to run tools (like search, calculator, DB query) → return `"continue"`.

Graph

```
graph = StateGraph(AgentState)

graph.add_node("our_agent", model_call)

#Tool Node
tool_node = ToolNode(tools=tools)
graph.add_node("tools", tool_node) #Contains multiple tools

graph.set_entry_point("our_agent")

graph.add_conditional_edges(
    "our_agent",
    should_continue,
    {
        "continue": "tools",
```

```

        "end" : END
    }
)

#Edge that goes back to our tool
graph.add_edge("tools", "our_agent")

app = graph.compile()

```

`graph.add_edge("tools", "our_agent")`: Edge that goes back to our tool.

- But we also wanted `tools` → `our_agent`

Stream

```

def print_stream(stream):
    for s in stream:
        message = s["messages"][-1]
        if isinstance(message, tuple):
            print(message)
        else:
            message.pretty_print()

```

for s in stream:

- This means: go through each **item** (`s`) inside `stream`.
- If `stream` is a list (or an iterable object), this loop will pick one item at a time.

Think of `stream` like a **queue of updates/messages** coming from an AI pipeline — you're checking them one by one.

message = s["messages"][-1]

- `s` is expected to be a dictionary (key-value pairs).
- Inside `s`, there is a key `"messages"`.
- `s["messages"]` is a list of messages.

- `[-1]` means "last item in the list".

`if isinstance(message, tuple):`

- `isinstance(x, tuple)` checks if `message` is a **tuple** (a fixed group of values like `(a, b)` in Python).
- Why? Because sometimes the last message might be stored as a **tuple** instead of a custom object.
- If `message` is not a tuple, it must be some kind of **object** (likely a `ChatMessage` or `AIMessage` from LangChain).
- **That object has a built-in method `.pretty_print()` which formats it nicely for display.**

```
inputs = {"messages": [("user", "Add 40 + 12")]}
print_stream(app.stream(inputs, stream_mode="values"))
```

```
===== Human Message =====
Add 40 + 12
===== Ai Message =====
Tool Calls:
  add (6gjdhgmks)
  Call ID: 6gjdhgmks
  Args:
    a: 40
    b: 12
===== Tool Message =====
Name: add

52
===== Ai Message =====

The result of the addition is 52.
```

- `.stream(inputs, stream_mode="values")` runs the pipeline but **instead of giving you the final result all at once**, it yields outputs **step by step** as a **stream** (like a

generator).

- Each yielded value (`s`) is a dictionary containing intermediate + final messages.

```
inputs = {"messages": [("user", "Add 40 + 12. add 5+7")]}  
print_stream(app.stream(inputs, stream_mode="values"))
```

```
===== Human Message =====  
  
Add 40 + 12. add 5+7  
  
===== Ai Message =====  
Tool Calls:  
  add (asx4mnnn7) 1  
Call ID: asx4mnnn7  
  Args:  
    a: 40  
    b: 12  
  add (fkxb2dq20) 2  
Call ID: fkxb2dq20  
  Args:  
    a: 5  
    b: 7  
  
===== Tool Message =====  
Name: add  
  
12  
  
===== Ai Message =====  
  
The results of the additions are 52 and 12.
```

Visualize

```
from IPython.display import Image, display  
display(Image(app.get_graph().draw_mermaid_png()))
```

